

Let  $x$  point to an element of a doubly linked list. Let  $L[x]$  and  $R[x]$  point to the predecessor and successor of that element.

Then the operations are:  $L/$ bracketleft,  $R/$ Bracketright,  $L/$  Bracket

$L[x]$  and  $R[x]$  no longer have their former semantic significance after  $x$  has been removed. This fact is, of course, obvious, once it has been pointed out.

Danger lurks when objects are allowed to point into a list from the

The element denoted by  $x$  has been deleted from its list; why would anybody want to put it back again? Well, I admit that updates to a data structure are usually intended to be permanent

interactive programs may need to revert to a former state when the user wants to undo an operation or a sequence of operations. Another typical

Backtracking, also called depth-

The idea of (2) was introduced in 1979 by Hitotumatu and Noshita. It makes Dijkstra's

Backtrack programming can be re-garded as the task of deciding how to narrow the search and at the same time to organize the data that controls the decisions. Floyd's elegant discussion of the connection between backtracking and nondeterministic algorithms includes a precise method for updating data structures before choosing

In simple situations we can simply maintain a stack that contains snapshots of the relevant state information at all ancestors of the current node in the search tree. But the task of copying the entire state at each level might take too much

1

Dijkstra's recursive procedure for the queens problem kept the current state in three global Boolean arrays. Hitotumatu and Noshita's program kept it in a doubly

When Dijkstra tentatively placed a queen, he changed one entry of each Boolean array from true to false; then he made the entry true again when backtracking. Hitotumatu and Noshita used (1) to remove a column and

The beauty of (2) is

General schemes for undoing assignments require us to record the identity of the left-hand side together with its previous value. But in this case only the single quantity  $x$  is needed

We can apply (1) and (2) repeatedly in complex data structures that involve large numbers of interacting doubly linked lists. The program logic that traverses those lists and decides what elements should be deleted can often be run in reverse

This process causes the pointer variables inside the global data structure to execute an exquisitely choreographed

Given a matrix of 0s and 1s, does it have a set of rows containing exactly one 1 in each column? We can think of the columns as elements of a universe, and the rows as subsets of the universe; then the problem is to cover the universe.

Dana Scott conducted one of the first experiments on backtracking programming in 1958. His program, written for the MANI

12 pentominoes are placed in

Scott was probably inspired by Golomb's paper [14] and some extensions reported by Martin Gardner. For example, one of the 65 solutions is shown in Figure

construct all possible rows representing a way to place a pentomino on the board. Each row contains a 1 in the column identifying the piece, and 1s in the columns identifying its positions.

Algorithm X is simply a statement of the obvious trial-and-error approach. Algorithm X solves all solutions to the exact cover problem.

If  $A$  is empty

Choose a row,

Include  $r$  in the partial solution

For each  $j$  such that  $A[r, j] =$

Repeat this algorithm recursively

The nondeterministic choice of  $r$  means that the algorithm essentially clones itself into subalgorithms. Each subalgorithm inherits the current matrix  $A$ , but reduces

The subalgorithms form a search tree in a natural way, with the original problem

Backtracking is the process of traversing the tree in pre-order, "depth first." Any systematic rule for choosing column  $c$  in this procedure will provide all solutions. But some rules work much better than others.

Scott realized that piece X has only 3 essentially different positions, namely centered at 23, 24, and 33. The full set of 8

$8 \times 65 = 520$  solutions is easily obtained by rotation and reflection.

Golomb and Baumert suggested choosing, at each stage of a backtrack procedure, a subproblem that leads to the fewest branches. In the case of an exact cover problem, this

Search trees for Scott's pentomino problem then have only 10,421 nodes (X at 23);

Rows of the matrix are doubly linked as circular lists via the `LandR` fields. Each column list also includes a special data object called its list header.

The list headers are part of a larger object called a column object. The size is the number of 1s in the column, and the name is a symbolic identifier for printing the answers.

The `LandR` fields of the list headers link together all columns that still need to be covered. This circular list also includes a

The 0-1 matrix of (3) would be represented by the objects shown in Figure 2. For example, if we name the columns A, B, C, D, E, F, and G, this is the matrix. The Clinks are not shown because they would clutter up the picture.

Cover column  $c$  (see below)

For each  $r$ , while  $r \neq c$ , set  $Ok[r]$ ; for each  $j$ , while  $j \neq r$ , uncover column  $j$  (see

Uncover column  $c$  (see below) and return

The operation of printing the current solution is easy. We successively print the rows containing data object  $O_0, O_1, \dots, O_{k-1}$ .

5

Four-way-linked representation of the exact

To choose a column object  $c$ , we could simply set  $c = R[h]$ ; this is the leftmost uncovered column. Or if we want to

The operation of covering column  $c$  is more interesting: It removes  $c$  from the header list and removes all rows in  $c$ 's own list from the other column lists they

Set  $L \leftarrow \text{bracketleftbig}$

For each  $i$ , while  $i \neq c$ , set  $S \leftarrow \text{bracketleftbig} C[i] \text{bracketrightbig}$ . For each  $j$ , set  $R \leftarrow \text{bracketleftbig}$

Operation (1) is used to remove objects in both the horizontal and

For each  $i = U[c], U/\text{bracketleftbig } U[c]/\text{bracketrightbig}$ , while  $i/\text{negationslash} = c$ , we get to the point of this whole algorithm, the operation of uncovering a given column  $c$ . Here is where

$R/\text{bracketleftbig}$

6

Uncovering takes place in precisely the reverse order of the covering operation. Notice that (2) undoes (1)

Figure 2 shows what happens when search (0) is applied to the data of (3) Figure 3 shows the asymmetry of the links that now appear in column D.

Search (1) will cover column B, and there will be no 1s left in column E. Continuing search (0), when  $r$  points to the A element of row (A, D, G), we also cover columns D and G. So search (2) will find nothing.

7

The links after columns D and G in Figure 3 have

If the  $S^*$ elds are ignored in the choice of  $c$ , or if the shortest column is chosen at each step, the solution will be found.

Readers who play through the action of this algorithm on some examples

The running time of algorithm DL X is essentially proportional to the number of times it applies operation (1) to remove an object from a list. Efficiency considerations. When algorithm X is implemented in terms of dancing links, let's call it algorithm DLX.

The search tree for one

Study larger examples before drawing any general conclusions

A backtrack program usually spends most of its time on only a few levels of the search tree. Figure 5 shows the search tree for the case  $X = 23$  of Dana Scott's pentomino problem using the heuristic.

More than half of the nodes lie on levels  $\leq 8$ . Extra work on the lower levels has reduced the need for hard work at the higher levels.

Each update involves about 14 memory accesses when the heuristic is used, and about 8 when  $S$  is ignored.

The heuristic is even more effective in larger problems, because it tends to reduce the total number of nodes by a factor that is exponential in the number of levels. The cost of applying it grows only linearly.

The Sheuristic is good in large trees but not so good in small ones. I tried a hybrid scheme that uses the Sheuristic at low levels but not at high levels. This experiment was, however, unsuccessful. Therefore I decided to retain the heuristic at all levels of algorithm DLX.

10

My trusty old SPARC station 2 computer, vintage 1992, is able to perform approximately 0.39 mega-updates per second when working on large problems. The 120 MHz Pentium I computer that Stanford computer science faculty were given in 1996 did 1.21 mega-updates per second.

Dancing links can be used to solve complex problems with simple geometric structure. The technique of dancing links is actually a step backward from Scott's 40-year-old method.

The task of packing the set of pentominoes into a  $6 \times 10$  rectangle is more difficult than Scott's  $8 \times 8 \times 2 \times 2$  problem. The backtrack tree for the  $6 \times 10$  problem is larger and there are 2339 essentially different solutions [21]. In this case we limit the Xpentomino to the upper left

John G. Fletcher needed only ten minutes to solve the  $6 \times 10$  problem on an IBM

The 7094 had a clock rate of 0.7 MHz, and it could access two 36-bit words in a single clockcycle. Fletcher's program required only about  $600 \times 700,000 / 28,320,810$  clock cycles.

Many people have written about polyomino problems, including distinguished mathematic

92 solutions, 14,352,556 nodes, 1,764,631,796 updates. 100 solutions, 10,258,180

Algorithm DLX will branch on the ways to place a cell if some piece is difficult to place. It knows no difference, because pieces and cells are simply columns of the given input matrix.

Algorithm DLX begins to outperform other pentomino-placing procedures in problems where the search tree has many levels. For example, let's consider the problem

The Multum produced an answer after more than an hour, but she remained uncertain whether other solutions were possible. Now, with the dancing links approach described above, we can obtain several solutions almost instantly. The solutions fall into four classes,

12

In the late 1950s, T. H. O'Beirne introduced a pleasant variation on polyominoes by substituting triangles for squares.

The twelve hexiamonds were independently discovered by J. E. Ree

The  $6 \times 6$  rhombus can be tiled by the twelve hexiamonds in exactly 156 ways. Figure 7 shows one such arrangement, together with some arrow dissections.

13

O'Sheare was particularly fascinated by the fact that seven of the twelve hexiamonds have different shapes when they are flipped over. In November of 1959, after three months of trials, he found a solution; and

A 19-level backtrack tree with many possibilities at each level makes an excellent test case for the dancing links approach to covering the puzzle. I broke the general case into seven subcases, depending on the distance of the hexagon piece from the center. The total number of updates performed was 134,425,768,494.

My goal was not only to count the solutions, but also to count arrangements that were as symmetrical as possible. The overall hexagon has 156 internal edges, and the 19 one-sided hexiamonds have 96 internal non-edges.

A solution to the hexiamond problem is maximally symmetric if it has the highest horizontal or vertical symmetry score.

Four of the eight solutions shown in Figure 8 have a horizontal symmetry score of 32. John Conway found these solutions by hand in 1964 and conjectured that they were symmetric overall.

14

Figure 8 shows the solutions to O'Sheare's hexiamond hexagon problem. The small hexagon is at various distances from the center of the large one.

15

There are 46 ways to pack the one-sided pentominoes in a  $3 \times 30$  rectangle. Figure 9 shows a maximally symmetric example (which isn't really very symmetrical)

46 solutions, 605,440 nodes, 190,311,7

I set out to count the solutions to the  $9 \times 10$  problem. I soon found that the task would be hopeless, unless I invented a much better algorithm. The Monte Carlo estimation procedure of [24] suggests that about 19 quadrillion updates will be needed, with 64 trillion nodes in the search

I do, however, have a conjecture about the

A failed experiment. Special arguments based on "coloring" often give important insights into tiling problems.

For example, it is well known that if we remove two cells from opposite corners of a chessboard, there is no way to cover the remaining 62 cells with dominoes.

Algorithm DLX makes 4,780,846 updates before concluding that there is no

The cells of the hexiamond-hexagon problem can be colored black and white in a similar fashion. All triangles that point left are black, say, and all that point right are white. The remaining four, namely the \*sphinx\* and the \*

I expected the subproblems to run up to 16 times as fast as the original problem. I expected the extra information about impossible correlations of piece placement to help algorithm DLX make intelligent choices.

The overall problem had 6675 solutions and required 8,976,245,858 updates. The six subproblems turned out to have respectively 955, 1208, 1164, 1106, 1272, and 970 solutions.

Brian Barwell considered making tetrasticks from line segments or sticks. He called the resulting objects polysticks.

There are 2 disticks, 5 tristicks, and 16 tetrastick.

Barwell proved that the sixteen tetrasticks cannot be assembled into any symmetrical shape. But by leaving out any one of the tetrastick that have an excess of horizontal or vertical line segments, he found ways to create a 5x5 square.

Such puzzles are quite difficult to do by hand, and he had found only five solutions at the time he wrote his paper.

Generalized problem asks for a set of rows that covers every primary column exactly once and every secondary column at most once.

The tetrastick problem of Figure 11(c) can be set up as a general

Figure 11. Filling a 5x5 grid with 15 of the 16 tetrasticks; we must leave out either the H, the J, the L, the N, or the Y.

W, X, Y, Z representing the fifteen tetrasticks (excluding L) Columns Hxy representing the horizontal segments  $(x, y) \rightarrow (x+1, y)$ , and Vxy representing the vertical segments. Secondary columns Ixy represent interior junction points  $(x, y)$ , for  $0 <$

Polysticks are not supposed to

The common interior point I33 means that these rows cross each other. For example, the two rows corresponding to the Figure 11(c) covers only the interior points I14, I21, I

We can solve the generalized cover problem by using almost the same algorithm as before. The only difference is that we initialize the data structure by making a circular list of the column headers for the primary columns only. The header

A generalized cover problem can be converted to an equivalent exact cover problem. But we are better off working with the generalized problem, because the

There are ten one-sided welded tetrasticks if we add the mirror images of the unsymmetrical pieces. Only three solutions are possible, including the two perfectly symmetric solutions shown.

One-sided welded tetrasticks

There are 14 one-sided unwelded tetrasticks, and I thought they would surely fit into a 5x5 grid in a similar way. But this turned out to be impossible. The reason is that four of the six pieces J, J\*, L, L\*

Figure 14 shows

19

I also tried unsuccessfully to pack all 25 of the one-sided tetrasticks into the Aztec diamond pattern of Figure 15. I see no way to prove that a

The 4 queens problem is just the task of covering eight primary columns (R0, R1, R2, R3, F0, F1, F2, F3) corresponding to ranks and files. The problem is actually a special case of the generalized cover problem in the previous section.

When we apply algorithm DLX to this generalized cover problem, it behaves quite differently from the traditional algorithms for the Nqueens problem. It branches sometimes on different ways to occupy a rank of the chessboard and sometimes on files of a chessboard.

Central positions rule out more possibilities for later placements. We gain efficiency by paying attention to the order in which primary columns are considered.

Figure 16(a) shows an empty board with 8 possible ways to occupy each rank and each file. Placing a queen in R2 and F3 after Figure 16(d) makes it impossible to cover F2. Backtracking will occur even though only four queens have been tentatively placed.

Figure 16. Solving the 8 queens problem by treating ranks and files symmetrically.

21

The order in which header nodes are linked together at the start of algorithm DLX can have a significant effect on the running time. For example, the search tree has 312,512,659 nodes and requires 5,801,583,789 up-dates, if



Algorithm DLX solved small cases of the Nqueens problem using organ-pipe order. The advantage of mixing rows with columns becomes evident as the number of solutions increases.

Special methods are known for counting the number of solutions to the Nqueens

Algorithm DLX is an effective way to enumerate all solutions to such problems. On small cases it is nearly as fast as algorithms that have been tuned to solve particular classes of problems.

The ordering heuristic is used to tackle larger and larger cases all the time

In this paper I have used the exact cover problem to illustrate the versatility of dancing links. I could have chosen many other backtrack applications in which the same ideas apply. For example, the approach works nicely with the Waltz filtering algorithm [36]

I wish to thank Sol Golomb, Richard Guy, and Gene Freuder for the help they generously gave me as I was preparing this paper. Maggie McLoughlin

"I profoundly thank Tomas Rokicki, who provided the new computer on

My names for the tetrasticks are slightly different from those originally proposed by Barwell. I prefer to use the letters J, R, and U for the pieces he called U, J, and C respectively.

The implementation of algorithm DLX that I used when preparing this paper is the dance.w on webpage [http://www-](http://www-23)

This puzzle, which is available from [www.puzzletts.com](http://www.puzzletts.com), actually has only 83 solutions. It carries a Chinese title, "Dragon's Intelligence Pro't System."

[1] Brian R. Barwell

Elwyn R. Berlekamp, John H. Conway,

Max Black, Critical Thinking (Eng

Ole-Johan Dahl, Edsger W. Dijkstra, and C. R. Ho

N. G. de Bruijn, personal communication (9 September 1999): "... it was almost my first activity in programming that I got all 2

"I could speed the matter up by having a very long program, and that one was generated by

Henry Ernest Dudeney, "74" The broken chess

A program to solve the pentomino problem by the recursive

Robert W. Floyd, \*N

Martin Gardner, \*Mathematical games: More about complex dominoes

Michael R. Garey and David S

Solomon W. Golomb, \*Check

Solomon W. Golomb, Polyomino

Solomon W. Golomb and Leonard D.

24

Richard K. Guy, \*Some mathematical recreations,\*

Richard K. Guy, \*O\*Beirne\*s Hexiamond,

Robert M. Haralick and Gordon L. Elliott, \*Increasing t

Packing a square with Y-pentominoes

C. B. and Jennifer Haselgrove, \*A

Hiroshi Hitotumatu and Kohei Noshita, \*A

George P. Jelliss, \*

Donald E. Knuth, \*Estimating the

Donald E. Knuth, T

Jean Meeus, \*Some polyom

N. Metropolis and J. Worlton, \*

Pell\*s equation in two popular problems

T. H. O\*Beirne, \*Puzzles and Paradoxes 44: Pentominoes and hexiamonds,\* NewScientist 12(1961), 316\*317.

Hexiamond is a type of diamond

J. E. Reeve and J

Igor Rivin, Ilan Vardi, and Paul

Dana S. Scott,\*Programming a combinatorial puzzle,\* Technical Report No.1 (Princeton, New Jersey: Princeton

University Department of

P. J. Torbijn

David Waltz, \*Understanding line drawings of scenes with shadows,\* in

Bernhard Wierwille and Jacques Ha

Alfred Wassermann of Universität Bayreuth covered the Aztec diamond of Figure 15 with one-sided tetrasticks. The

107 possible solutions, which

Many of these turn out to be more symmetric than the one in Figure 10